

# Importing Libraries

```
In [1]: import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
```

# Importing Datasets

```
In [2]: df=pd.read_csv(r"C:\Users\user\Downloads\madrid_2008.csv")
df
```

		<b>date</b>	<b>BEN</b>	<b>CO</b>	<b>EBE</b>	<b>MXY</b>	<b>NMHC</b>	<b>NO<sub>2</sub></b>	<b>NOx</b>	<b>O<sub>XY</sub></b>	<b>O<sub>3</sub></b>	<b>PM10</b>	<b>PM2</b>
0	01-06-2008 01:00		NaN	0.47	NaN	NaN	NaN	83.089996	120.699997	NaN	16.990000	16.889999	10.4
1	01-06-2008 01:00		NaN	0.59	NaN	NaN	NaN	94.820000	130.399994	NaN	17.469999	19.040001	NaN
2	01-06-2008 01:00		NaN	0.55	NaN	NaN	NaN	75.919998	104.599998	NaN	13.470000	20.270000	NaN
3	01-06-2008 01:00		NaN	0.36	NaN	NaN	NaN	61.029999	66.559998	NaN	23.110001	10.850000	NaN
4	01-06-2008 01:00		1.68	0.80	1.70	3.01	0.30	105.199997	214.899994	1.61	12.120000	37.160000	21.9
...	...	...	...	...	...	...	...	...	...	...	...	...	...
226387	01-11-2008 00:00		0.48	0.30	0.57	1.00	0.31	13.050000	14.160000	0.91	57.400002	5.450000	5.1
226388	01-11-2008 00:00		NaN	0.30	NaN	NaN	NaN	41.880001	48.500000	NaN	35.830002	15.020000	NaN
226389	01-11-2008 00:00		0.25	NaN	0.56	NaN	0.11	83.610001	102.199997	NaN	14.130000	17.540001	13.9

	date	BEN	CO	EBE	MXY	NMHC	NO_2	NOx	OXY	O_3	PM10	PM2
226390	01-11-2008 00:00	0.54	NaN	2.70	NaN	0.18	70.639999	81.860001	NaN	NaN	11.910000	NaN
226391	01-11-2008 00:00	0.75	0.36	1.20	2.75	0.16	58.240002	74.239998	1.64	31.910000	12.690000	11.4

226392 rows × 17 columns

## Data Cleaning and Data Preprocessing

In [3]: `df=df.dropna()`In [4]: `df.columns`Out[4]: `Index(['date', 'BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3', 'PM10', 'PM25', 'PXY', 'SO_2', 'TCH', 'TOL', 'station'], dtype='object')`In [5]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 25631 entries, 4 to 226391
Data columns (total 17 columns):
 #   Column    Non-Null Count  Dtype  
--- 
 0   date      25631 non-null   object 
 1   BEN       25631 non-null   float64
 2   CO        25631 non-null   float64
 3   EBE       25631 non-null   float64
 4   MXY       25631 non-null   float64
 5   NMHC      25631 non-null   float64
 6   NO_2      25631 non-null   float64
 7   NOx       25631 non-null   float64
 8   OXY       25631 non-null   float64
 9   O_3        25631 non-null   float64
 10  PM10      25631 non-null   float64
 11  PM25      25631 non-null   float64
 12  PXY       25631 non-null   float64
 13  SO_2       25631 non-null   float64
 14  TCH       25631 non-null   float64
 15  TOL       25631 non-null   float64
 16  station    25631 non-null   int64  
dtypes: float64(15), int64(1), object(1)
memory usage: 3.5+ MB
```

In [6]: `data=df[['CO', 'station']]`  
`data`

Out[6]:

	CO	station
<b>4</b>	0.80	28079006
<b>21</b>	0.37	28079024
<b>25</b>	0.39	28079099
<b>30</b>	0.51	28079006
<b>47</b>	0.39	28079024
...	...	...
<b>226362</b>	0.35	28079024
<b>226366</b>	0.46	28079099
<b>226371</b>	0.53	28079006
<b>226387</b>	0.30	28079024
<b>226391</b>	0.36	28079099

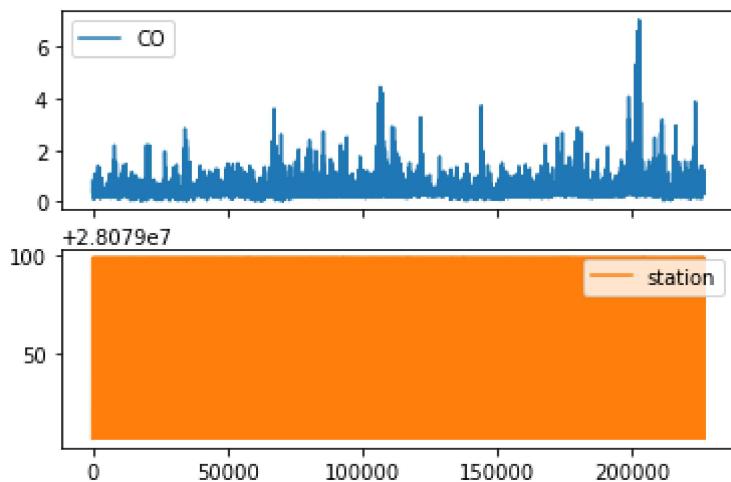
25631 rows × 2 columns

## Line chart

In [7]:

data.plot.line(subplots=True)

Out[7]: array([&lt;AxesSubplot:&gt;, &lt;AxesSubplot:&gt;], dtype=object)

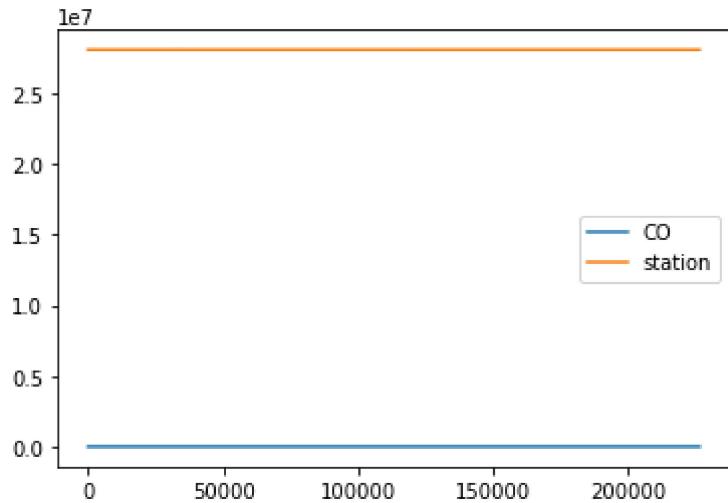


## Line chart

In [8]:

data.plot.line()

Out[8]: &lt;AxesSubplot:&gt;

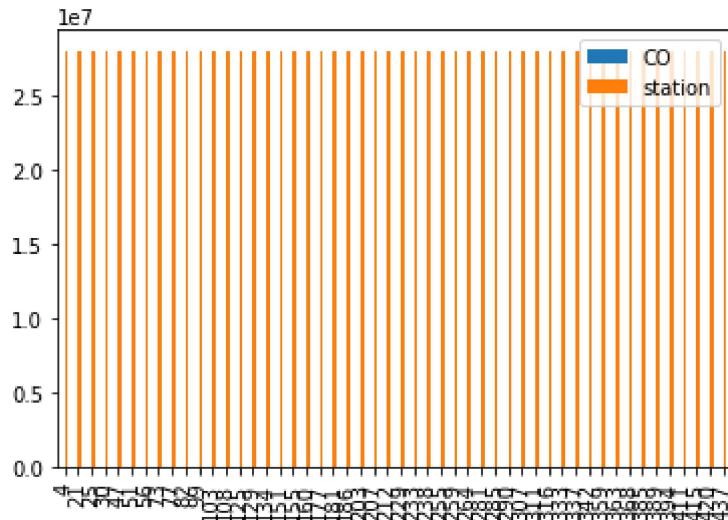


## Bar chart

```
In [9]: b=data[0:50]
```

```
In [10]: b.plot.bar()
```

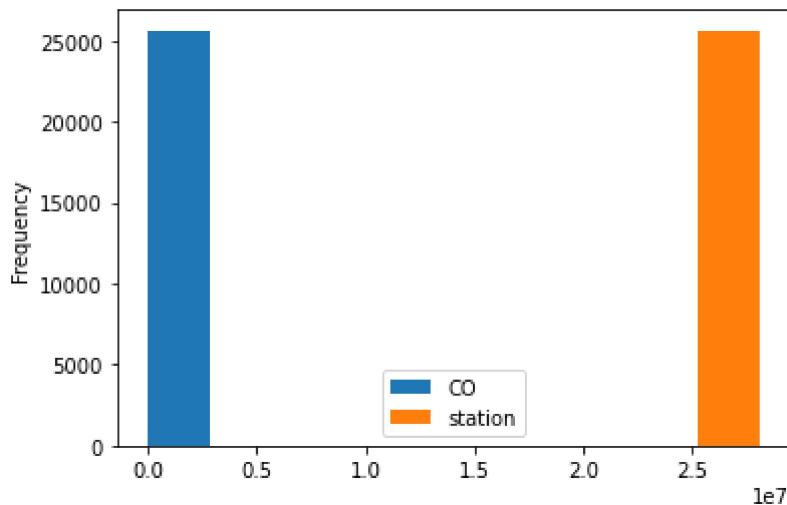
```
Out[10]: <AxesSubplot:>
```



## Histogram

```
In [11]: data.plot.hist()
```

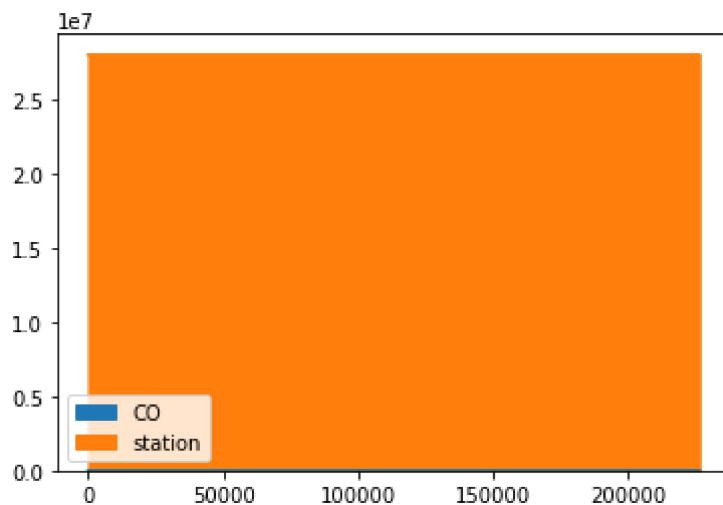
```
Out[11]: <AxesSubplot:ylabel='Frequency'>
```



## Area chart

In [12]: `data.plot.area()`

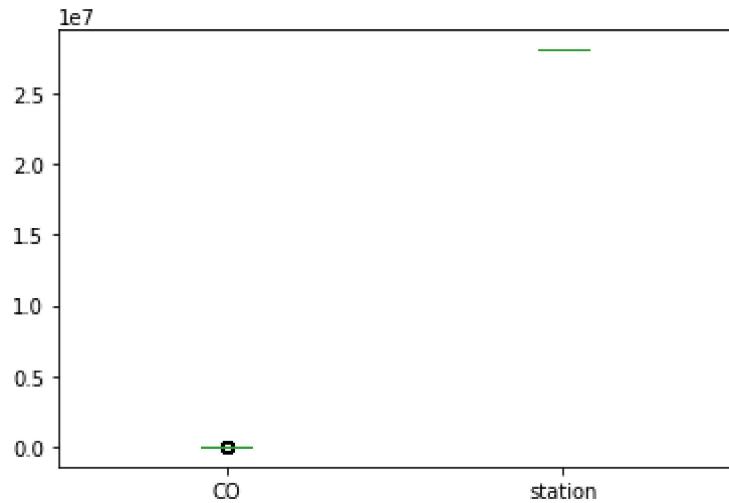
Out[12]: <AxesSubplot:>



## Box chart

In [13]: `data.plot.box()`

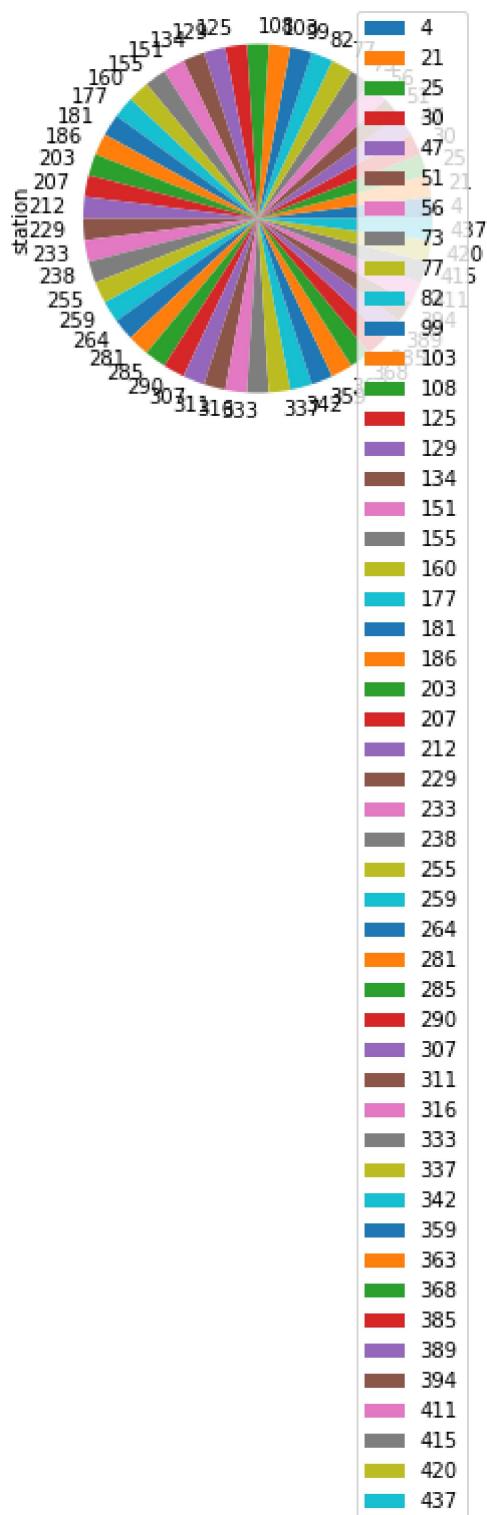
Out[13]: <AxesSubplot:>



## Pie chart

```
In [14]: b.plot.pie(y='station' )
```

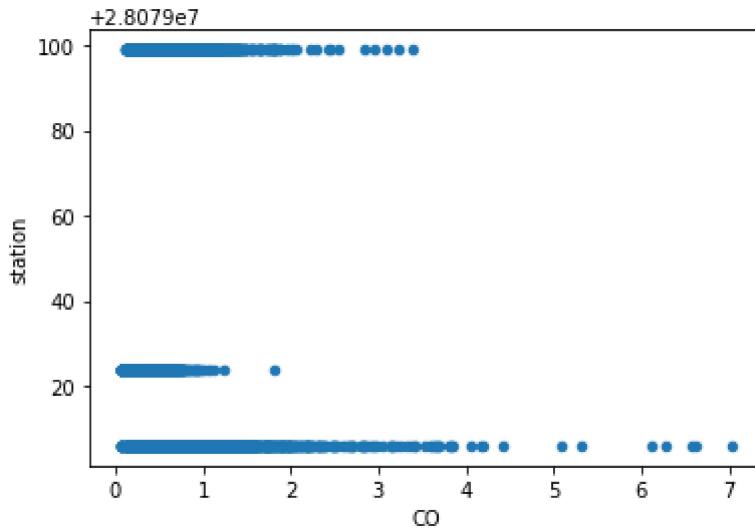
```
Out[14]: <AxesSubplot:ylabel='station'>
```



## Scatter chart

```
In [15]: data.plot.scatter(x='CO' ,y='station')
```

```
Out[15]: <AxesSubplot:xlabel='CO', ylabel='station'>
```



In [16]:

`df.info()`

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 25631 entries, 4 to 226391
Data columns (total 17 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   date        25631 non-null   object 
 1   BEN          25631 non-null   float64
 2   CO           25631 non-null   float64
 3   EBE          25631 non-null   float64
 4   MXY          25631 non-null   float64
 5   NMHC         25631 non-null   float64
 6   NO_2          25631 non-null   float64
 7   NOX          25631 non-null   float64
 8   OXY          25631 non-null   float64
 9   O_3           25631 non-null   float64
 10  PM10         25631 non-null   float64
 11  PM25         25631 non-null   float64
 12  PXY          25631 non-null   float64
 13  SO_2          25631 non-null   float64
 14  TCH          25631 non-null   float64
 15  TOL          25631 non-null   float64
 16  station       25631 non-null   int64  
dtypes: float64(15), int64(1), object(1)
memory usage: 3.5+ MB
```

In [17]:

`df.describe()`

Out[17]:

	BEN	CO	EBE	MXY	NMHC	NO_2	NO
count	25631.000000	25631.000000	25631.000000	25631.000000	25631.000000	25631.000000	25631.000000
mean	1.090541	0.440632	1.352355	2.446045	0.213323	54.225261	98.007731
std	1.146461	0.317853	1.118191	2.390023	0.123409	38.164647	101.448231
min	0.100000	0.060000	0.170000	0.240000	0.000000	0.240000	2.110000
25%	0.430000	0.260000	0.740000	1.000000	0.130000	25.719999	32.635000
50%	0.750000	0.350000	1.000000	1.620000	0.190000	48.000000	71.110000

	BEN	CO	EBE	MXY	NMHC	NO_2	NOx
75%	1.320000	0.510000	1.580000	3.105000	0.270000	74.924999	131.550000
max	27.230000	7.030000	26.740000	55.889999	1.760000	554.900024	2004.000000

In [18]:

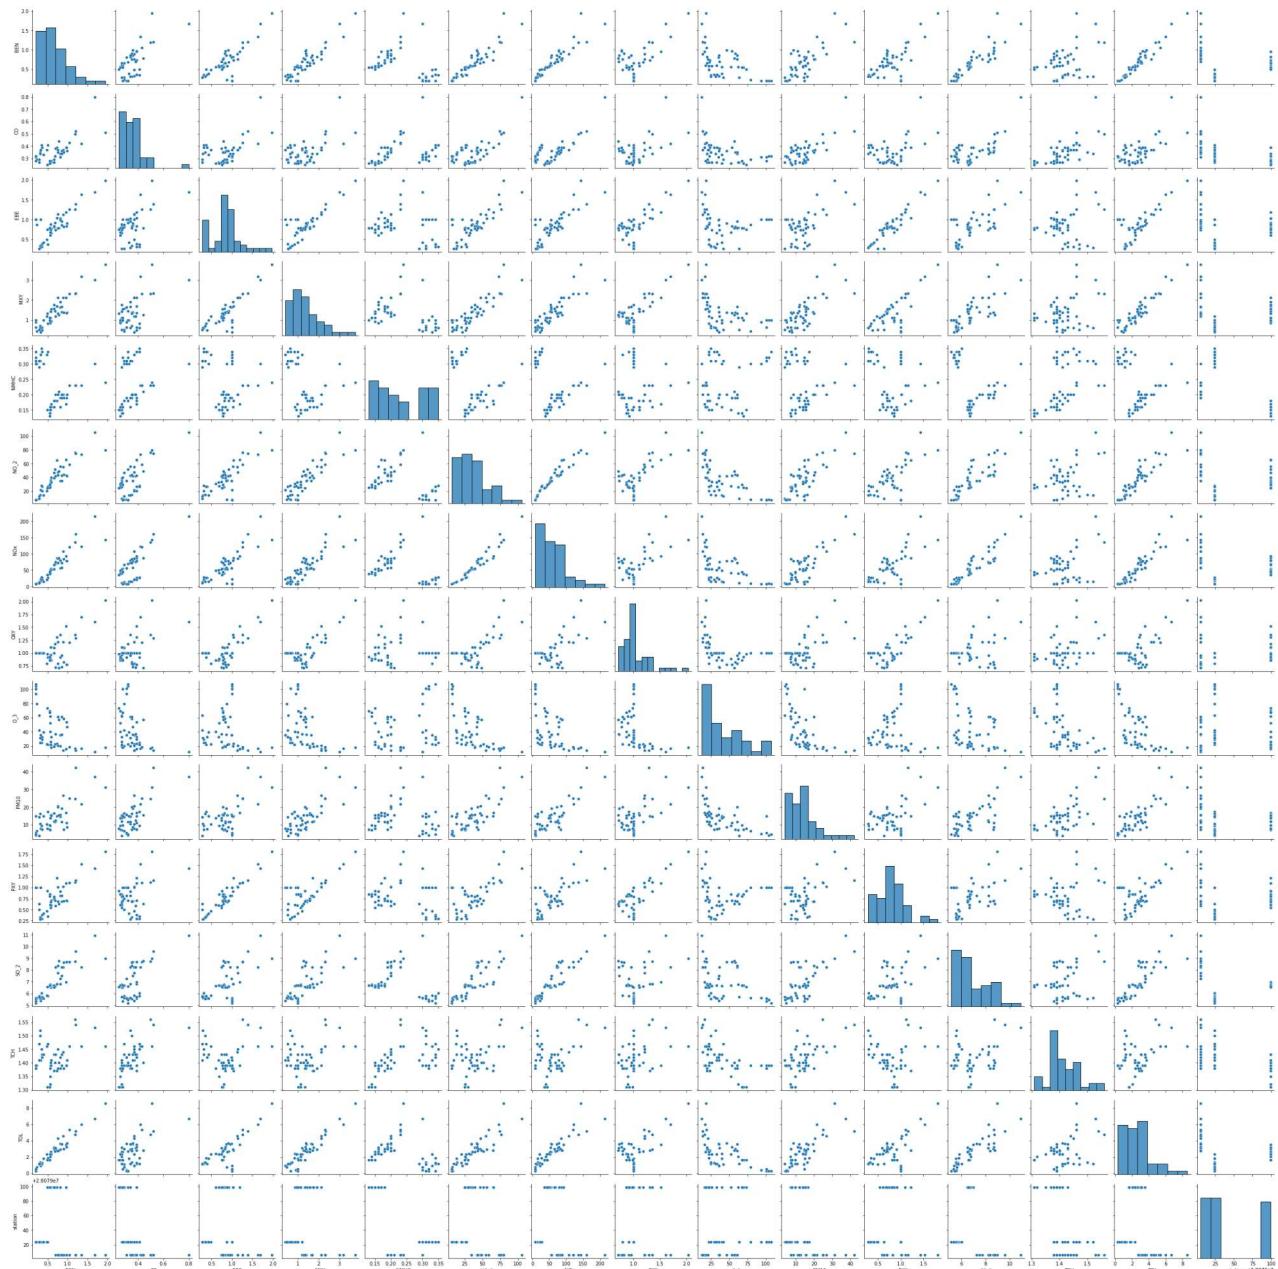
```
df1=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3',
       'PM10', 'PXY', 'SO_2', 'TCH', 'TOL', 'station']]
```

## EDA AND VISUALIZATION

In [19]:

```
sns.pairplot(df1[0:50])
```

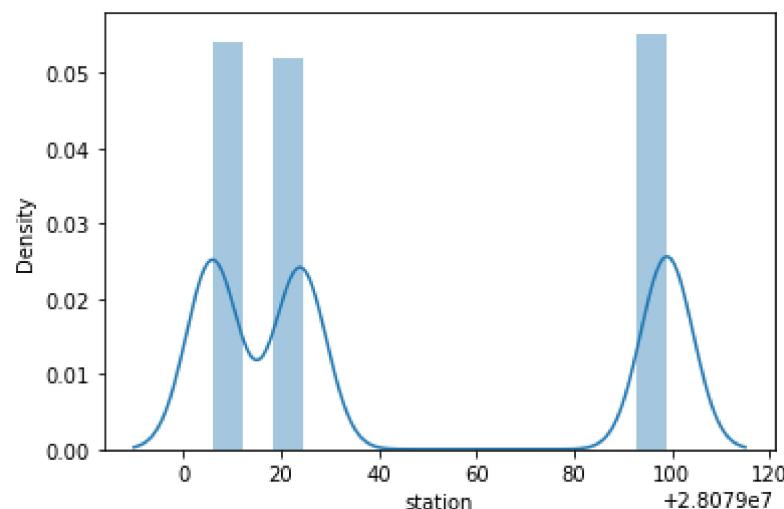
Out[19]:



In [20]: `sns.distplot(df1['station'])`

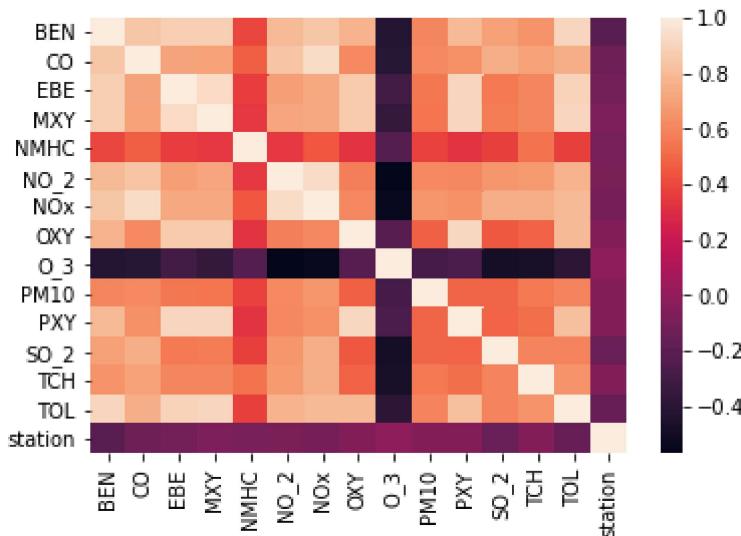
```
C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning:
`distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).
  warnings.warn(msg, FutureWarning)
```

Out[20]: <AxesSubplot:xlabel='station', ylabel='Density'>



In [21]: `sns.heatmap(df1.corr())`

Out[21]: <AxesSubplot:>



## TO TRAIN THE MODEL AND MODEL BUILDING

In [22]: `x=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3', 'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]  
y=df['station']`

```
In [23]: from sklearn.model_selection import train_test_split  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.3)
```

## Linear Regression

```
In [24]: from sklearn.linear_model import LinearRegression  
lr=LinearRegression()  
lr.fit(x_train,y_train)
```

```
Out[24]: LinearRegression()
```

```
In [25]: lr.intercept_
```

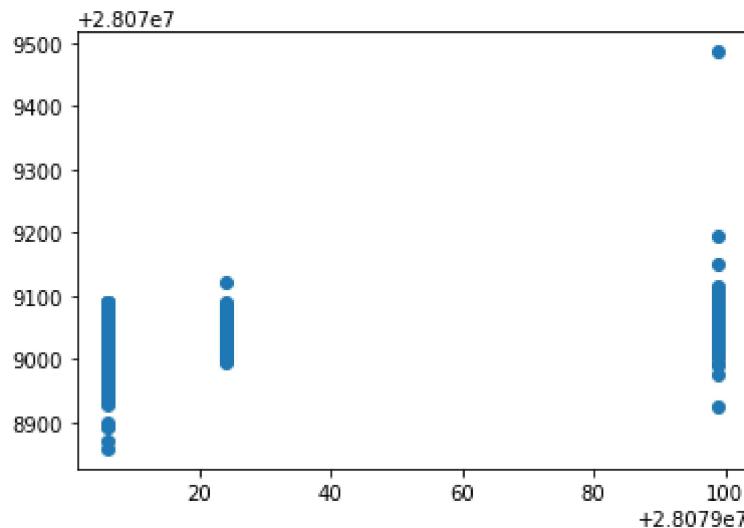
```
Out[25]: 28079036.397904124
```

```
In [26]: coeff=pd.DataFrame(lr.coef_,x.columns,columns=['Co-efficient'])  
coeff
```

	Co-efficient
BEN	-26.371935
CO	-1.853514
EBE	-2.407173
MXY	7.499778
NMHC	-25.948440
NO_2	-0.054696
NOx	0.140469
OXY	4.161625
O_3	-0.137308
PM10	0.132925
PXY	4.062823
SO_2	-0.566902
TCH	15.972244
TOL	-1.881657

```
In [27]: prediction =lr.predict(x_test)  
plt.scatter(y_test,prediction)
```

```
Out[27]: <matplotlib.collections.PathCollection at 0x2dee228e850>
```



## ACCURACY

```
In [28]: lr.score(x_test,y_test)
```

```
Out[28]: 0.13796631568894835
```

```
In [29]: lr.score(x_train,y_train)
```

```
Out[29]: 0.14563020243366387
```

## Ridge and Lasso

```
In [30]: from sklearn.linear_model import Ridge,Lasso
```

```
In [31]: rr=Ridge(alpha=10)  
rr.fit(x_train,y_train)
```

```
Out[31]: Ridge(alpha=10)
```

## Accuracy(Ridge)

```
In [32]: rr.score(x_test,y_test)
```

```
Out[32]: 0.13788369114445032
```

```
In [33]: rr.score(x_train,y_train)
```

```
Out[33]: 0.14560850699579742
```

```
In [34]: la=Lasso(alpha=10)
la.fit(x_train,y_train)
```

Out[34]: Lasso(alpha=10)

```
In [35]: la.score(x_train,y_train)
```

Out[35]: 0.040880310514757046

## Accuracy(Lasso)

```
In [36]: la.score(x_test,y_test)
```

Out[36]: 0.041939499058839425

## Elastic Net regression

```
In [37]: from sklearn.linear_model import ElasticNet
en=ElasticNet()
en.fit(x_train,y_train)
```

Out[37]: ElasticNet()

```
In [38]: en.coef_
```

Out[38]: array([-4.59605114, -0. , -0. , 3.36119425, -0. ,
 0.04586469, 0.03105838, 1.57283069, -0.15599659, 0.13430945,
 1.68310687, -0.95294493, 0. , -2.59604792])

```
In [39]: en.intercept_
```

Out[39]: 28079057.189837225

```
In [40]: prediction=en.predict(x_test)
```

```
In [41]: en.score(x_test,y_test)
```

Out[41]: 0.09581264928877309

## Evaluation Metrics

```
In [42]: from sklearn import metrics
print(metrics.mean_absolute_error(y_test,prediction))
```

```
print(metrics.mean_squared_error(y_test,prediction))
print(np.sqrt(metrics.mean_squared_error(y_test,prediction)))
```

35.7801596720135  
1494.585665856179  
38.65987151887832

## Logistic Regression

In [43]: `from sklearn.linear_model import LogisticRegression`

In [44]: `feature_matrix=df[['BEN', 'CO', 'EBE', 'MXY', 'NMHC', 'NO_2', 'NOx', 'OXY', 'O_3', 'PM10', 'PXY', 'SO_2', 'TCH', 'TOL']]
target_vector=df[ 'station']`

In [45]: `feature_matrix.shape`

Out[45]: (25631, 14)

In [46]: `target_vector.shape`

Out[46]: (25631,)

In [47]: `from sklearn.preprocessing import StandardScaler`

In [48]: `fs=StandardScaler().fit_transform(feature_matrix)`

In [49]: `logr=LogisticRegression(max_iter=10000)
logr.fit(fs,target_vector)`

Out[49]: `LogisticRegression(max_iter=10000)`

In [50]: `observation=[[1,2,3,4,5,6,7,8,9,10,11,12,13,14]]`

In [51]: `prediction=logr.predict(observation)
print(prediction)`

[28079099]

In [52]: `logr.classes_`

Out[52]: `array([28079006, 28079024, 28079099], dtype=int64)`

In [53]: `logr.score(fs,target_vector)`

```
Out[53]: 0.7942335453162186
```

```
In [54]: logr.predict_proba(observation)[0][0]
```

```
Out[54]: 8.280697732723258e-09
```

```
In [55]: logr.predict_proba(observation)
```

```
Out[55]: array([8.28069773e-09, 1.19826346e-13, 9.99999992e-01])
```

## Random Forest

```
In [56]: from sklearn.ensemble import RandomForestClassifier
```

```
In [57]: rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

```
Out[57]: RandomForestClassifier()
```

```
In [58]: parameters={'max_depth':[1,2,3,4,5],
                 'min_samples_leaf':[5,10,15,20,25],
                 'n_estimators':[10,20,30,40,50]
                }
```

```
In [59]: from sklearn.model_selection import GridSearchCV
grid_search =GridSearchCV(estimator=rfc,param_grid=parameters, cv=2,scoring="accuracy")
grid_search.fit(x_train,y_train)
```

```
Out[59]: GridSearchCV(cv=2, estimator=RandomForestClassifier(),
                      param_grid={'max_depth': [1, 2, 3, 4, 5],
                                  'min_samples_leaf': [5, 10, 15, 20, 25],
                                  'n_estimators': [10, 20, 30, 40, 50]},
                      scoring='accuracy')
```

```
In [60]: grid_search.best_score_
```

```
Out[60]: 0.8554711657916186
```

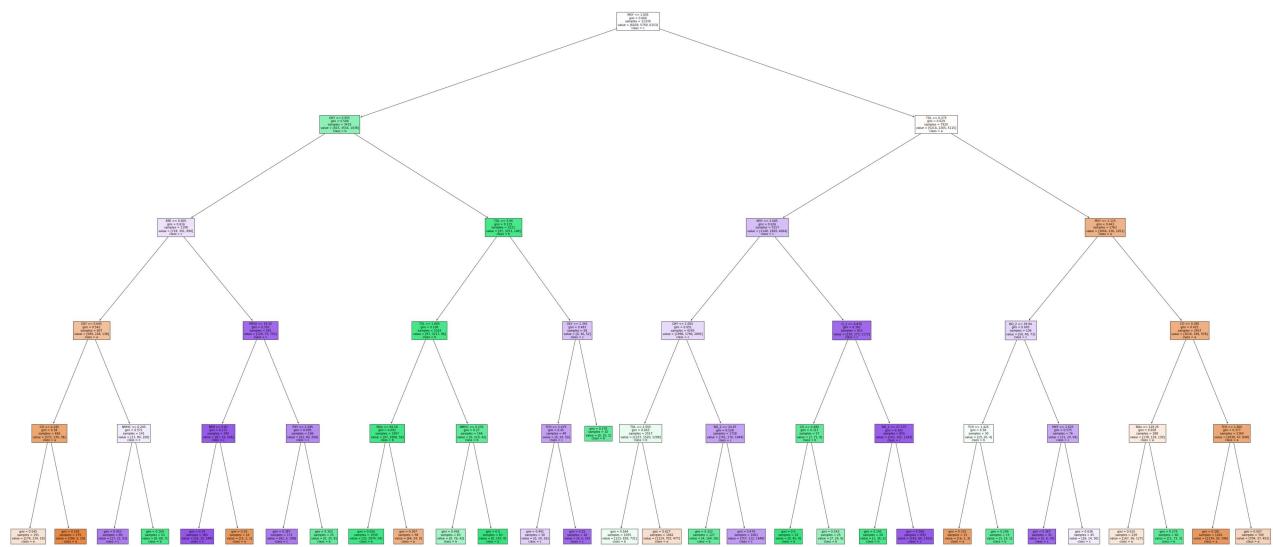
```
In [61]: rfc_best=grid_search.best_estimator_
```

```
In [62]: from sklearn.tree import plot_tree
```

```
plt.figure(figsize=(80,40))
plot_tree(rfc_best.estimators_[5],feature_names=x.columns,class_names=['a','b','c','d'])
```

```
Out[62]: [Text(2222.700000000003, 1993.2, 'MXY <= 1.005\ngini = 0.666\nsamples = 11339\nvalue = [6029, 5759, 6153]\nnclass = c'),
Text(1171.800000000002, 1630.800000000002, 'OXY <= 0.955\ngini = 0.508\nsamples = 3419\nvalue = [815, 3554, 1038]\nnclass = b'),
Text(595.2, 1268.4, 'EBE <= 0.605\ngini = 0.616\nsamples = 1198\nvalue = [718, 301, 890]\nnclass = c'),
Text(297.6, 906.0, 'OXY <= 0.645\ngini = 0.542\nsamples = 607\nvalue = [589, 228, 138]\nnclass = a'),
Text(148.8, 543.5999999999999, 'CO <= 0.245\ngini = 0.38\nsamples = 466\nvalue = [572, 139, 38]\nnclass = a'),
Text(74.4, 181.1999999999982, 'gini = 0.545\nsamples = 191\nvalue = [176, 134, 19]\nnclass = a'),
Text(223.200000000002, 181.1999999999982, 'gini = 0.109\nsamples = 275\nvalue = [396, 5, 19]\nnclass = a'),
Text(446.400000000003, 543.5999999999999, 'NMHC <= 0.205\ngini = 0.571\nsamples = 141\nvalue = [17, 89, 100]\nnclass = c'),
Text(372.0, 181.1999999999982, 'gini = 0.453\nsamples = 89\nvalue = [17, 21, 93]\nnclass = c'),
Text(520.800000000001, 181.1999999999982, 'gini = 0.169\nsamples = 52\nvalue = [0, 68, 7]\nnclass = b'),
Text(892.800000000001, 906.0, 'PM10 <= 19.32\ngini = 0.355\nsamples = 591\nvalue = [129, 73, 752]\nnclass = c'),
Text(744.0, 543.5999999999999, 'BEN <= 0.81\ngini = 0.272\nsamples = 395\nvalue = [67, 33, 546]\nnclass = c'),
Text(669.6, 181.1999999999982, 'gini = 0.24\nsamples = 385\nvalue = [52, 32, 544]\nnclass = c'),
Text(818.400000000001, 181.1999999999982, 'gini = 0.29\nsamples = 10\nvalue = [15, 1, 2]\nnclass = a'),
Text(1041.600000000001, 543.5999999999999, 'PXY <= 1.295\ngini = 0.495\nsamples = 196\nvalue = [62, 40, 206]\nnclass = c'),
Text(967.2, 181.1999999999982, 'gini = 0.387\nsamples = 171\nvalue = [62, 5, 198]\nnclass = c'),
Text(1116.0, 181.1999999999982, 'gini = 0.303\nsamples = 25\nvalue = [0, 35, 8]\nnclass = b'),
Text(1748.4, 1268.4, 'TOL <= 3.04\ngini = 0.133\nsamples = 2221\nvalue = [97, 3253, 148]\nnclass = b'),
Text(1488.0, 906.0, 'TOL <= 1.805\ngini = 0.108\nsamples = 2163\nvalue = [97, 3217, 96]\nnclass = b'),
Text(1339.2, 543.5999999999999, 'NOx <= 94.18\ngini = 0.092\nsamples = 1997\nvalue = [97, 2998, 54]\nnclass = b'),
Text(1264.800000000002, 181.1999999999982, 'gini = 0.056\nsamples = 1939\nvalue = [33, 2974, 54]\nnclass = b'),
Text(1413.600000000001, 181.1999999999982, 'gini = 0.397\nsamples = 58\nvalue = [64, 24, 0]\nnclass = a'),
Text(1636.800000000002, 543.5999999999999, 'NMHC <= 0.235\ngini = 0.27\nsamples = 166\nvalue = [0, 219, 42]\nnclass = b'),
Text(1562.4, 181.1999999999982, 'gini = 0.458\nsamples = 83\nvalue = [0, 76, 42]\nnclass = b'),
Text(1711.2, 181.1999999999982, 'gini = 0.0\nsamples = 83\nvalue = [0, 143, 0]\nnclass = b'),
Text(2008.800000000002, 906.0, 'OXY <= 1.345\ngini = 0.483\nsamples = 58\nvalue = [0, 36, 52]\nnclass = c'),
Text(1934.4, 543.5999999999999, 'TCH <= 1.475\ngini = 0.45\nsamples = 48\nvalue = [0, 26, 50]\nnclass = c'),
Text(1860.000000000002, 181.1999999999982, 'gini = 0.491\nsamples = 30\nvalue = [0, 20, 26]\nnclass = c'),
Text(2008.800000000002, 181.1999999999982, 'gini = 0.32\nsamples = 18\nvalue = [0, 6, 24]\nnclass = c'),
Text(2083.200000000003, 543.5999999999999, 'gini = 0.278\nsamples = 10\nvalue = [0, 10, 2]\nnclass = b'),
Text(3273.600000000004, 1630.800000000002, 'TOL <= 6.275\ngini = 0.629\nsamples = 7920\nvalue = [5214, 2205, 5115]\nnclass = a'),
Text(2678.4, 1268.4, 'MXY <= 2.685\ngini = 0.626\nsamples = 5157\nvalue = [2148, 1969, 4064]\nnclass = c'),
Text(2380.8, 906.0, 'OXY <= 1.005\ngini = 0.651\nsamples = 4245\nvalue = [1998, 1796, 2
```

```
892]\nclass = c'),  
    Text(2232.0, 543.599999999999, 'TOL <= 2.555\ngini = 0.663\nsamples = 2517\nvalue = [1  
237, 1520, 1208]\nclass = b'),  
    Text(2157.600000000004, 181.1999999999982, 'gini = 0.564\nsamples = 1055\nvalue = [12  
3, 818, 731]\nclass = b'),  
    Text(2306.4, 181.1999999999982, 'gini = 0.627\nsamples = 1462\nvalue = [1114, 702, 47  
7]\nclass = a'),  
    Text(2529.600000000004, 543.599999999999, 'NO_2 <= 24.07\ngini = 0.528\nsamples = 172  
8\nvalue = [761, 276, 1684]\nclass = c'),  
    Text(2455.200000000003, 181.1999999999982, 'gini = 0.322\nsamples = 127\nvalue = [4,  
164, 36]\nclass = b'),  
    Text(2604.0, 181.1999999999982, 'gini = 0.479\nsamples = 1601\nvalue = [757, 112, 164  
8]\nclass = c'),  
    Text(2976.0, 906.0, 'O_3 <= 6.875\ngini = 0.362\nsamples = 912\nvalue = [150, 173, 117  
2]\nclass = c'),  
    Text(2827.200000000003, 543.599999999999, 'CO <= 0.485\ngini = 0.317\nsamples = 57\nv  
alue = [7, 71, 9]\nclass = b'),  
    Text(2752.8, 181.1999999999982, 'gini = 0.0\nsamples = 32\nvalue = [0, 45, 0]\nclass =  
b'),  
    Text(2901.600000000004, 181.1999999999982, 'gini = 0.543\nsamples = 25\nvalue = [7, 2  
6, 9]\nclass = b'),  
    Text(3124.8, 543.599999999999, 'NO_2 <= 27.175\ngini = 0.302\nsamples = 855\nvalue =  
[143, 102, 1163]\nclass = c'),  
    Text(3050.4, 181.1999999999982, 'gini = 0.145\nsamples = 20\nvalue = [1, 36, 2]\nclass  
= b'),  
    Text(3199.200000000003, 181.1999999999982, 'gini = 0.268\nsamples = 835\nvalue = [14  
2, 66, 1161]\nclass = c'),  
    Text(3868.8, 1268.4, 'MXY <= 2.125\ngini = 0.443\nsamples = 2763\nvalue = [3066, 236, 1  
051]\nclass = a'),  
    Text(3571.200000000003, 906.0, 'NO_2 <= 39.94\ngini = 0.645\nsamples = 106\nvalue = [5  
0, 40, 73]\nclass = c'),  
    Text(3422.4, 543.599999999999, 'TCH <= 1.425\ngini = 0.58\nsamples = 30\nvalue = [19,  
20, 4]\nclass = b'),  
    Text(3348.000000000005, 181.1999999999982, 'gini = 0.335\nsamples = 15\nvalue = [16,  
1, 3]\nclass = a'),  
    Text(3496.8, 181.1999999999982, 'gini = 0.299\nsamples = 15\nvalue = [3, 19, 1]\nclass  
= b'),  
    Text(3720.000000000005, 543.599999999999, 'MXY <= 1.635\ngini = 0.575\nsamples = 76\nv  
alue = [31, 20, 69]\nclass = c'),  
    Text(3645.600000000004, 181.1999999999982, 'gini = 0.367\nsamples = 31\nvalue = [5,  
6, 39]\nclass = c'),  
    Text(3794.4, 181.1999999999982, 'gini = 0.638\nsamples = 45\nvalue = [26, 14, 30]\ncla  
ss = c'),  
    Text(4166.400000000001, 906.0, 'CO <= 0.385\ngini = 0.425\nsamples = 2657\nvalue = [301  
6, 196, 978]\nclass = a'),  
    Text(4017.600000000004, 543.599999999999, 'NOx <= 119.25\ngini = 0.658\nsamples = 288  
\nvalue = [178, 129, 130]\nclass = a'),  
    Text(3943.200000000003, 181.1999999999982, 'gini = 0.615\nsamples = 228\nvalue = [16  
7, 56, 127]\nclass = a'),  
    Text(4092.000000000005, 181.1999999999982, 'gini = 0.279\nsamples = 60\nvalue = [11,  
73, 3]\nclass = b'),  
    Text(4315.200000000001, 543.599999999999, 'TCH <= 1.605\ngini = 0.377\nsamples = 2369  
\nvalue = [2838, 67, 848]\nclass = a'),  
    Text(4240.8, 181.1999999999982, 'gini = 0.281\nsamples = 1620\nvalue = [2134, 30, 396]  
\nclass = a'),  
    Text(4389.6, 181.1999999999982, 'gini = 0.507\nsamples = 749\nvalue = [704, 37, 452]\ncla  
ss = a')]
```



# Conclusion

## Accuracy

### linear regression

```
In [63]: lr.score(x_test,y_test)
```

```
Out[63]: 0.13796631568894835
```

### Ridge regression

```
In [64]: rr.score(x_test,y_test)
```

```
Out[64]: 0.13788369114445032
```

### Lasso regression

```
In [65]: la.score(x_test,y_test)
```

```
Out[65]: 0.041939499058839425
```

### Elastic net regression

```
In [66]: en.score(x_test,y_test)
```

```
Out[66]: 0.09581264928877309
```

## Logistic regression

```
In [67]: logr.score(fs,target_vector)
```

```
Out[67]: 0.7942335453162186
```

## Random forest

```
In [68]: grid_search.best_score_
```

```
Out[68]: 0.8554711657916186
```

**Accuracy for random forest is higher so it is the best fit model**